

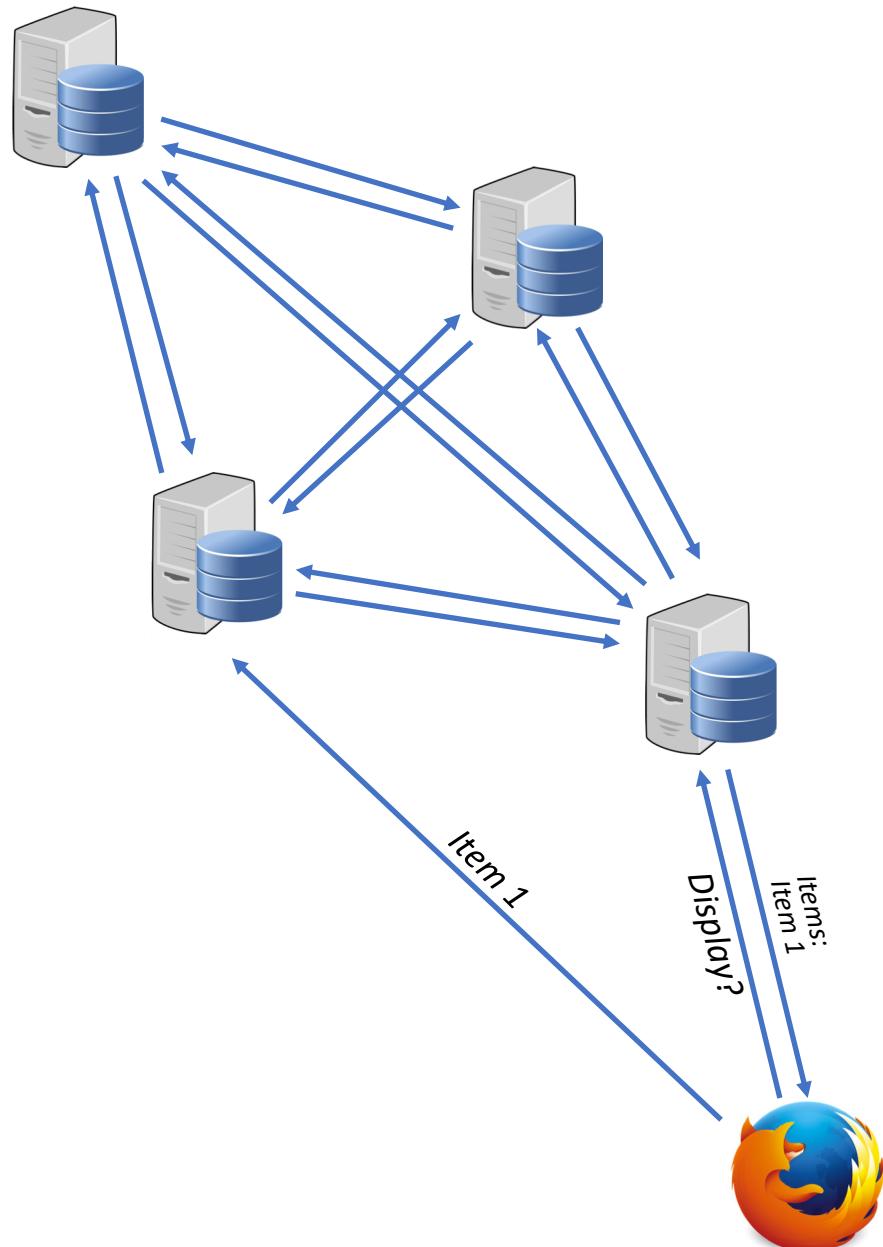
# Lab 2

## Leader Election and Centralized Solution



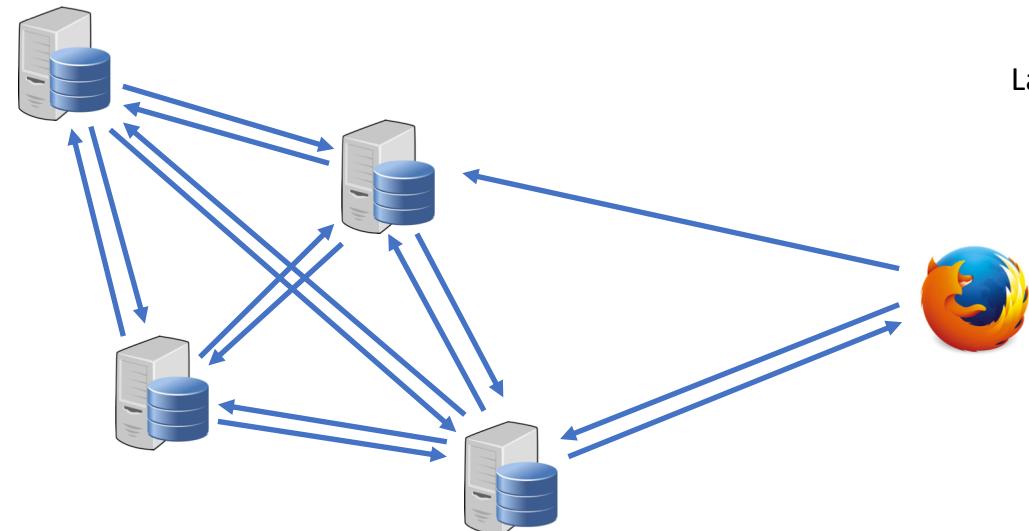
# Wrap-up lab 1: An Overly Simple Distributed Blackboard?

- Our blackboard *is* distributed
- But might get inconsistent!
- Consistent blackboard:
  - All entries are the same
  - Messages shown in the same order
  - Order is somehow associated with a notion of time
- How can we have a consistent blackboard?
  - Let's try a centralized solution!

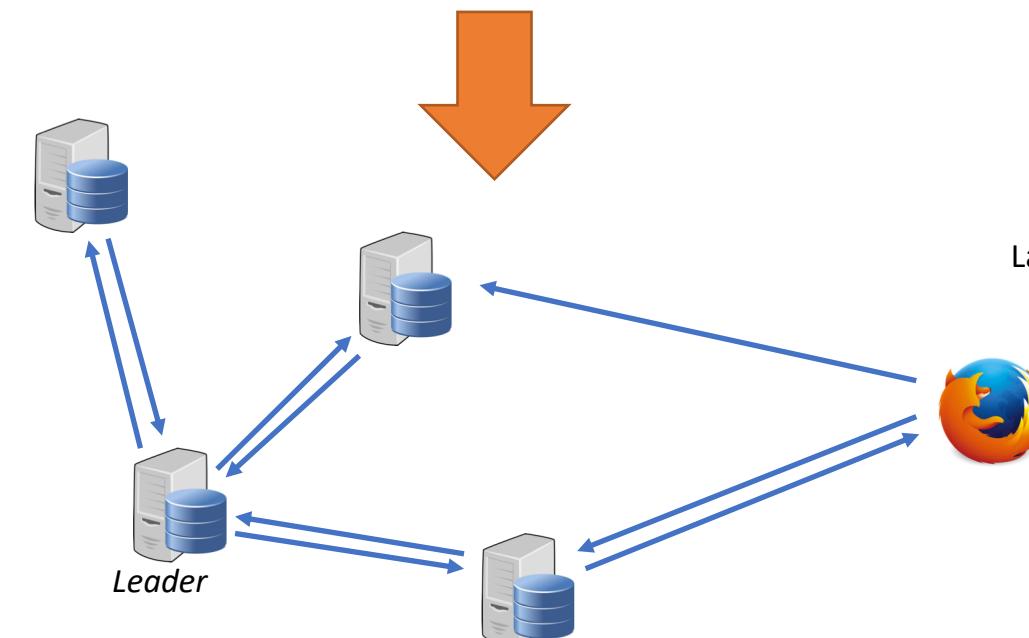


# Idea: a centralized solution

- Each post is sent to the leader
- Leader decides what to do and distributes it to the network
- Leader should be able to handle messages from multiple nodes and come up with a unique ordering
- Who will be leader?
  - We must execute a Leader Election (LE)
  - See lecture slides for different LE algorithms



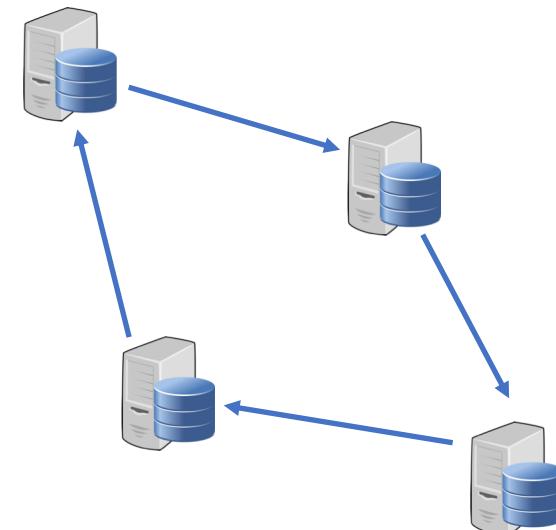
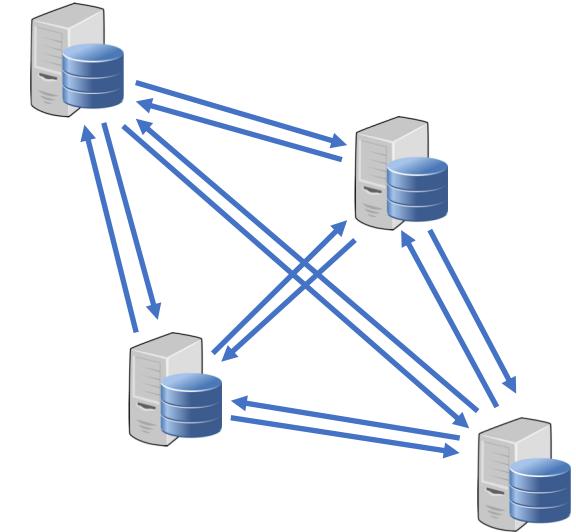
Lab 1



Lab 2

# Leader Election

- Which LE algorithm should I use?
  - Bully algorithm?
  - Ring algorithm?
  - Your own LE algorithm?
  - You decide!
- Which election attribute?
  - Plenty of choices: random number, IP or MAC address, some uuid...
  - **This lab: random number, generated when the server boots up**
    - **Some servers might have the same number! You can use the server ID (or IP address) as tie breaker**



# Lab 2 -Tasks

What you NEED to implement



# Task 1 – Leader Election

- Implement a *Leader Election* in your distributed blackboard
  - **All servers should start an election at the beginning!**
  - **LE attribute: random number selected at bootup + server ID as tie breaker**
- Hint: LE should be started at the beginning, but requires the server to be started to work properly. Start LE as a thread, with a delay to let the server start. Use `do_parallel_task_after_delay` just before `bottle.run`, with a few seconds delay
- Display on each server their election attribute, as well as their status (leader, slave)

## Task 2 – Centralized Blackboard

- Implement your *Centralized Blackboard*
  - Add, Modify and Delete should work
- Show the strong consistency of your solution
  - Order is the same on all servers
  - Concurrent modify do not break consistency

## Task 3 – Dynamic reelection after failures

- During runtime, kill your leader (by closing the terminal running bottle, or shutting down the link within mininet)
- Your system must detect that the leader has failed, and start a new election
  - You can reuse the random number generated at bootup
- What happens if the old leader is rebooted? (with loss of previous state-information)

## Task 4 – Argue for your choices

In a small text file uploaded with your code

- Discuss the solution cost of your LE
- Discuss whether your LE fulfills the properties of a correct LE algo
- Discuss the cost of adding messages to the board
- Briefly discuss pros&cons of centralized solutions

# Lab 2 - deliverables

What you **NEED** to give/show us



# Deliverables

- Your implementation
  - Including any script you used to test consistency!
  - Upload it on iLearn
- A video of your work
  - ~ 5 min (max 10)
  - Upload on the CAU-cloud
    - Everybody will be able to watch your video!
  - Should include (at least) Design – Demo – Evaluation
    - see Lab presentations for details

**Deadline:**  
**Wednesday 9<sup>th</sup>, December**

Good luck, and start coding!



# Additional slides

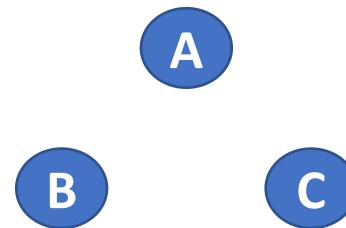


# Agenda

- Lab 2
- Solution cost for Lab 1
- Mininet Python API (optional)
- A small note on threads & CURL

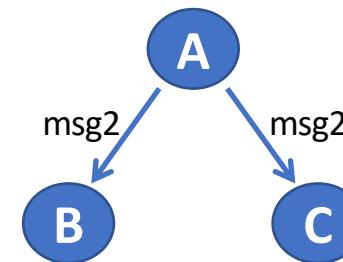
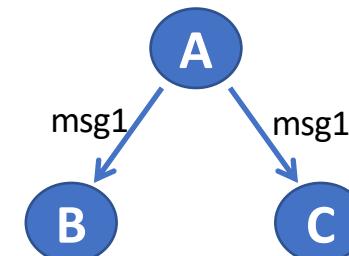
# Solution cost

- We care about the solution cost in terms of communication
- We can measure the cost of a solution (e.g. in Lab 1) in terms of
  - number of nodes to which a new post is propagated
  - payload: number of blackboard entries per message
- For example, consider the case of three nodes A, B and C, and the following events:
  - Event 1: User posts “msg1” to node A.
  - Event 2: User posts “msg2” to node A.



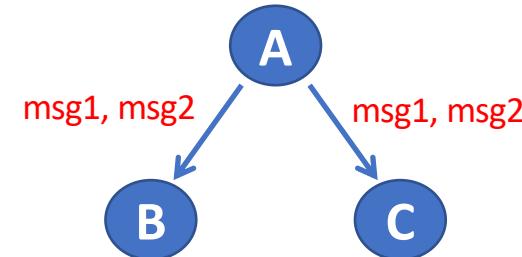
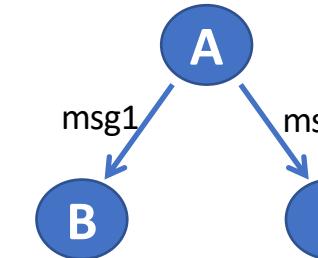
## Example: a good scenario

- Propagate only the new post.
- Upon Event 1,  
Vessel A sends “msg1” to vessels B and C
  - Payload for each message = 1
  - Overall cost = 2
- Upon Event 2,  
Vessel A sends “msg2” to vessels B and C
  - Payload for each message = 1
  - Overall cost = 2
- Overall cost per post: #nodes-1



## Example: a costly scenario

- Propagate the whole blackboard
- Upon Event 1, Vessel A sends “msg1” to vessels B and C
  - Payload for each message = 1
  - Overall cost = 2
- Upon Event 2, Vessel A sends “msg1, msg2” to vessels B and C
  - Payload for each message = 2
  - Overall cost = 4
- Overall cost per post:  $b(n-1)$ 
  - board size \* (#nodes-1)



## Cost of Lab2

- We want you do the same simple communication cost analysis for Lab2:
  - First for the leader election part.
    - What is the cost of the whole leader election?
  - Then for the centralized blackboard.
    - What is the cost of a new post?

# Agenda

- Lab 2
- Solution cost for Lab 1
- Mininet Python API (optional)
- A small note on threads & CURL

# Mininet Python API

- Staff that are not really required for the labs, but it is good to know ☺
- You can use a Python API to:
  - Set up your own topologies.
  - Configure link properties e.g. bandwidth, delay etc.
  - Run arbitrary commands on the hosts.
  - much more...
- ...all through python scripts.
- For example, *start\_topology.py* uses that API.

# Example 1: Building a custom topology (from the Mininet tutorial)

Defines a topology with n nodes  
connected to a single switch

```
class SingleSwitchTopo(Topo):  
  
    def build(self, n=2):  
        switch = self.addSwitch('s1')  
        # Python's range(N) generates 0..N-1  
        for h in range(n):  
            host = self.addHost('h%ss' % (h + 1))  
            self.addLink(host,  
switch)
```

# Example 1: Building a custom topology (from the Mininet tutorial)

Defines a topology with n nodes connected to a single switch

```
class SingleSwitchTopo(Topo):  
  
    def build(self, n=2):  
        switch = self.addSwitch('s1')  
        # Python's range(N) generates 0..N-1  
        for h in range(n):  
            host = self.addHost('h%ss' % (h + 1))  
            self.addLink(host,  
switch)
```

Inherit class *Topo* and override method *build*

add a switch

create a host

link it to the switch

# Example 1: Building a custom topology (from the Mininet tutorial)

Defines a topology with n nodes connected to a single switch

```
class SingleSwitchTopo(Topo):  
  
    def build(self, n=2):  
        switch = self.addSwitch('s1')  
        # Python's range(N) generates 0..N-1  
        for h in range(n):  
            host = self.addHost('h%ss' % (h + 1))  
            self.addLink(host,  
switch)
```

Starts Mininet with the specified topology

```
def simpleTest():  
  
    topo = SingleSwitchTopo(n=4)  
    net = Mininet(topo)  
    net.start()  
    print "Dumping host connections"  
    dumpNodeConnections(net.hosts)  
    print "Testing network connectivity"  
    net.pingAll()  
    net.stop()
```

## Example 2: Configuring link parameters (from the Mininet tutorial)

```
class SingleSwitchTopo(Topo):  
  
    def build(self, n=2):  
        switch = self.addSwitch('s1')  
        # Python's range(N) generates 0..N-1  
        for h in range(n):  
            host = self.addHost('h%0s' % (h + 1))  
            self.addLink( host, switch, bw=10, delay='5ms', loss=2 )
```

configure link properties

bandwidth in Mbs

delay in ms

loss rate %

## Example 2: Configuring link parameters (from the Mininet tutorial)

```
class SingleSwitchTopo(Topo):
```

```
    def build(self, n=2):
        switch = self.addSwitch('s1')
        # Python's range(N) generates 0..N-1
        for h in range(n):
            host = self.addHost('h%s' % (h + 1))
            self.addLink( host, switch, bw=10, delay='5ms', loss=2 )
```

Test the bandwidth  
between two nodes

or run any command on a  
node

```
h1, h4 = net.get( 'h1', 'h4' )
net.iperf( h1, h4 )
```

```
result = h1.cmd('ifconfig')
print result
```

# Mininet Python API

- If you want to know more:
  1. here is walkthrough of the API  
<https://github.com/mininet/mininet/wiki/Introduction-to-Mininet>
  2. repo with many example scripts  
<https://github.com/mininet/mininet/tree/master/examples>
  3. Lots of videos about Mininet, making custom topologies etc.  
<https://github.com/mininet/mininet/wiki/Videos>
  4. or look at the script that we gave you for lab 1 (lab1.py).

# Agenda

- Lab 2
- Solution cost for Lab 1
- Mininet Python API (optional)
- A small note on threads & CURL

# A small note about threads(1)

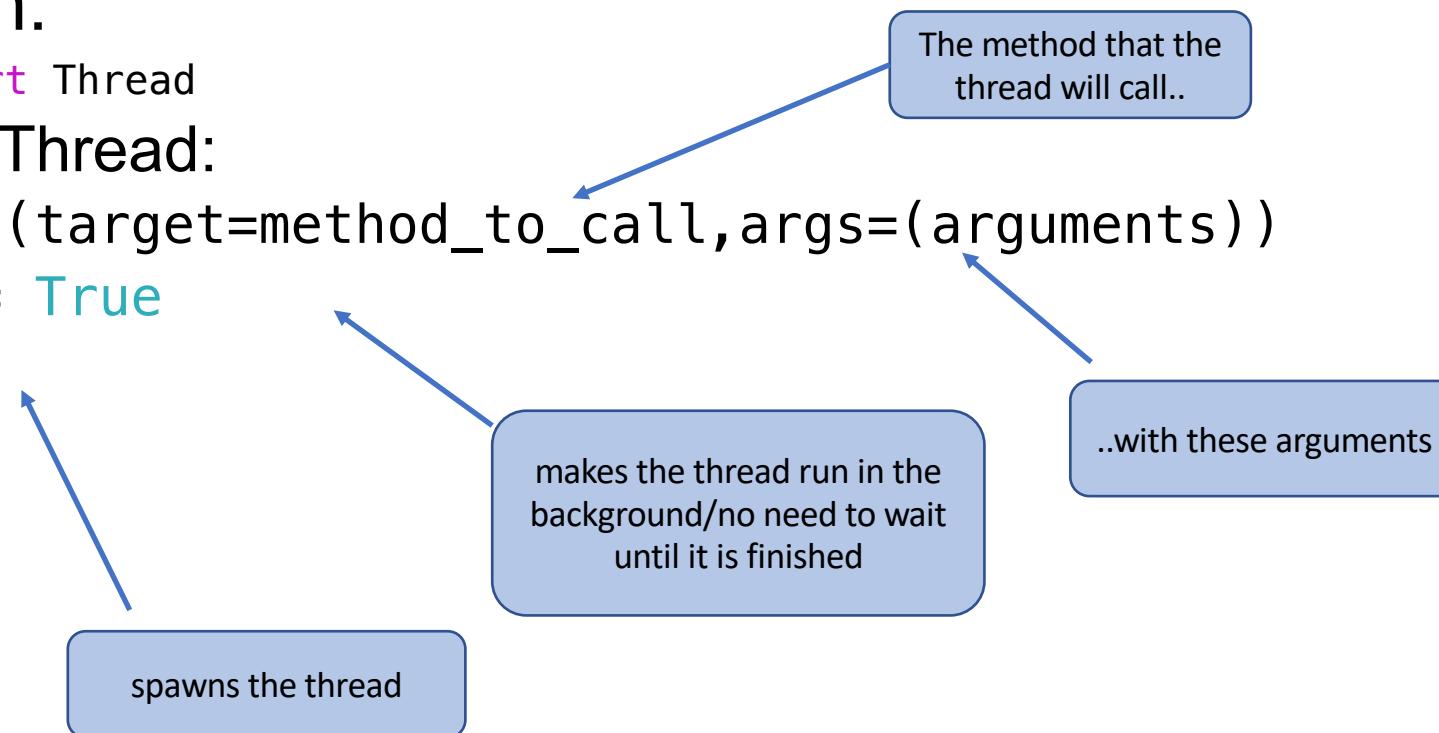
- At the *server.py* that we give you
  - Inside method *do\_POST*

```
thread = Thread(target=self.server.propagate_value_to_vessels,args=("action", "key", "value") )  
thread.daemon = True  
thread.start()
```

- The idea is to offload the propagation of messages to a different thread.
- So that *do\_POST* can return immediately.

# A small note about threads(2)

- Threads in python:
  - `from threading import Thread`
  - How to spawn a Thread:
    - 1) `t = Thread(target=method_to_call, args=(arguments))`
    - 2) `t.daemon = True`
    - 3) `t.start()`



## A small note about threads(3)

- No need to worry much about it in the labs.
- Use the code as it is when you want to propagate.

... but if you want to know more...

- the threading library:
  - <https://docs.python.org/3/library/threading.html>
- An interesting presentation (advanced):
  - <https://www.slideshare.net/dabeaz/an-introduction-to-python-concurrency>

# CURL

- There is more than Firefox
  - HTTP POST on the command line

```
curl -d 'entry=t'${i} -X 'POST' 'http://ip:port/entries'
```

done

This will **post** entry=t1, entry=t2, ..., entry=t20 to <http://ip:port/entries>

# General notes about the labs

- Mininet is used as an infrastructure to test your servers on complex topologies with many nodes.
- Servers are “plain” Python, that can run everywhere.
- You can test them locally, on your own machine.
  - Might need to change the default port and how to contact the neighbors.
  - E.g. Start two servers on localhost, one on port 8080 and one on 8081.
- Then switch to Mininet to test with a bigger topology.